



IBM Developer  
SKILLS NETWORK

# Winning Space Race with Data Science

Friederike Meyer  
06.10.2022



# Outline

---

- Executive Summary
- Introduction
- Methodology
- Results
- Conclusion
- Appendix

# Executive Summary

---

## Summary of methodologies

- Data Collection (REST API and Web Scraping)
- Data Wrangling
- Data Visualization and SQL
- Interactive Folium Map
- Dashboard
- Classification Models

## Summary of all results

- EDA Results
- Interactive Folium Maps
- Predictive Analytics

# Introduction

---

- SpaceX offers the Falcon 9 rocket at a cost of 62 million dollar. The company is able to offer this low price because they can reuse the first stage.
- We predict whether the Falcon 9 first stage will land successfully.



Section 1

# Methodology

# Methodology

---

## Executive Summary

- Data collection methodology:
  - Rest API
  - Web Scraping
- Perform data wrangling
  - Transformation of variables
- Perform exploratory data analysis (EDA) using visualization and SQL
- Perform interactive visual analytics using Folium and Plotly Dash
- Perform predictive analysis using classification models
  - Classification Models (Logistics Regression, Support Vector Machine, Decision Trees, K Nearest Neighbors)

# Data Collection

---

## Two Data Sources:

- SpaceX REST API
- Web scraping from Wikipedia

## Data Collection Steps:

- 1) SpaceX REST API to import complete data set on SpaceX Launches
- 2) Web scraping with BeautifulSoup from Wikipedia to import the data contained on the table on the Wikipedia page

# Data Collection – SpaceX API

- Data Collection with SpaceX REST API calls

Link to notebook:

[https://jupyter-dataplatform.cloud.ibm.com/analytics/notebooks/v2/e2f6ed1a-b996-42e5-8b8b-365f49dcc90b/view?access\\_token=2c8112bca28e0ebf2536623f2bbfbed4a16068e02673366715abdf14b282949f](https://jupyter-dataplatform.cloud.ibm.com/analytics/notebooks/v2/e2f6ed1a-b996-42e5-8b8b-365f49dcc90b/view?access_token=2c8112bca28e0ebf2536623f2bbfbed4a16068e02673366715abdf14b282949f)

## Flowchart of Key SpaceX API calls

```
In [31]: spacex_url="https://api.spacexdata.com/v4/launches/past"
```

```
In [32]: response = requests.get(spacex_url)
```

```
In [36]: # Use json_normalize meethod to convert the json result into a dataframe  
data=pd.json_normalize(response.json())
```

```
In [41]: # Call getBoosterVersion  
getBoosterVersion(data)
```

```
In [58]: # Hint data['BoosterVersion']!='Falcon 1'  
  
indexnames = launch_data.loc[launch_data['BoosterVersion']!='Falcon 9'].index  
launch_data.drop(indexnames, inplace = True)  
data_falcon9 = launch_data  
data_falcon9.head()
```



# Data Collection - Scraping

- Data Collection with WebScraping

Link to notebook:

[https://jp-tok.dataplatform.cloud.ibm.com/analytics/notebooks/v2/5a6fffb3-fa69-4462-884c-fd1b9211f862/view?access\\_token=4807bc84d6fdb1cd0acdda5609dd629afc682068d2333e9c61ea8339ebc17300](https://jp-tok.dataplatform.cloud.ibm.com/analytics/notebooks/v2/5a6fffb3-fa69-4462-884c-fd1b9211f862/view?access_token=4807bc84d6fdb1cd0acdda5609dd629afc682068d2333e9c61ea8339ebc17300)

## Flowchart of Key SpaceX API calls

```
In [31]: # use requests.get() method with the provided static_url
# assign the response to a object

data=requests.get(static_url)
```

Create a BeautifulSoup object from the HTML response

```
In [35]: # Use BeautifulSoup() to create a BeautifulSoup object from a response text content
soup=BeautifulSoup(data.content, 'html5lib')
```

Print the page title to verify if the BeautifulSoup object was created properly

```
In [36]: # Use soup.title attribute
soup.title
```

```
Out[36]: <title>List of Falcon 9 and Falcon Heavy launches - Wikipedia</title>
```

```
In [37]: # Use the find_all function in the BeautifulSoup object, with element type `table`
# Assign the result to a list called `html_tables`
html_tables=soup.find_all('table')
```

```
In [39]: column_names = []

# Apply find_all() function with `th` element on first_launch_table
# Iterate each th element and apply the provided extract_column_from_header() to get a column name
# Append the Non-empty column name ('if name is not None and len(name) > 0') into a list called column_names

elements=first_launch_table.find_all('th')
for element in elements:
    column_name=extract_column_from_header(element)
    if column_name!=None and len(column_name)>0:
        column_names.append(column_name)
```

Next Step: Creating Dictionary and Appending Data to Keys

# Data Wrangling

---

## Data Wrangling Steps:

- 1) Calculate the number of launches on each site
- 2) Calculate the number and occurrence of each orbit
- 3) Calculate the number and occurrence of mission outcomes per orbit type
- 4) Create a landing outcome label from the Outcome column

Link to Notebook:

[https://jp-tok.dataplatform.cloud.ibm.com/analytics/notebooks/v2/939d7293-f5a2-4b15-864d-d9d1d51c023f/view?access\\_token=d6cf9c2079a6430e94f5eba9c606c4d815a7693eec161021c58a1ee9935dbf34](https://jp-tok.dataplatform.cloud.ibm.com/analytics/notebooks/v2/939d7293-f5a2-4b15-864d-d9d1d51c023f/view?access_token=d6cf9c2079a6430e94f5eba9c606c4d815a7693eec161021c58a1ee9935dbf34)

# EDA with Data Visualization

---

## Overview of Plotted Charts:

- Scatterplot (for Continuous Variables)
- Bar Chart (for Categorical Variables)

Link to Notebook: [https://jp-tok.dataplatform.cloud.ibm.com/analytics/notebooks/v2/dcd67f60-d631-475f-8f2a-ce9a8031783d/view?access\\_token=ae11ae1ef7163e0fbf70c46d7d246a465093d6a733bd08857140f5b3f604cec2](https://jp-tok.dataplatform.cloud.ibm.com/analytics/notebooks/v2/dcd67f60-d631-475f-8f2a-ce9a8031783d/view?access_token=ae11ae1ef7163e0fbf70c46d7d246a465093d6a733bd08857140f5b3f604cec2)

# EDA with SQL

---

## EDA with SQL Steps:

- 1) Connect to Database
- 2) Display the names of the unique launch sites in the space mission
- 3) Display 5 records where launch sites begin with the string 'CCA'
- 4) Display the total payload mass carried by boosters launched by NASA (CRS)
- 5) Display average payload mass carried by booster version F9 v1.1
- 6) List the date when the first successful landing outcome in ground pad was achieved
- 7) List the names of the boosters which have success in drone ship and have payload mass greater than 4000 but less than 6000
- 8) List the total number of successful and failure mission outcomes
- 9) List the names of the booster\_versions which have carried the maximum payload mass. Use a subquery
- 10) List the failed landing\_outcomes in drone ship, their booster versions, and launch site names for in year 2015
- 11) Rank the count of landing outcomes (such as Failure (drone ship) or Success (ground pad)) between the date 2010-06-04 and 2017-03-20, in descending order

Link to Notebook: [https://jp-tok.dataplatform.cloud.ibm.com/analytics/notebooks/v2/0e958e7d-725f-4b75-8095-febc578ff7c6/view?access\\_token=80e5a8ee0a35a6778ba0615f3e7e2a7c6ac08d33c15f17ce2c33ae9ef8260ede](https://jp-tok.dataplatform.cloud.ibm.com/analytics/notebooks/v2/0e958e7d-725f-4b75-8095-febc578ff7c6/view?access_token=80e5a8ee0a35a6778ba0615f3e7e2a7c6ac08d33c15f17ce2c33ae9ef8260ede)

# Build an Interactive Map with Folium

---

## Overview of Map Items:

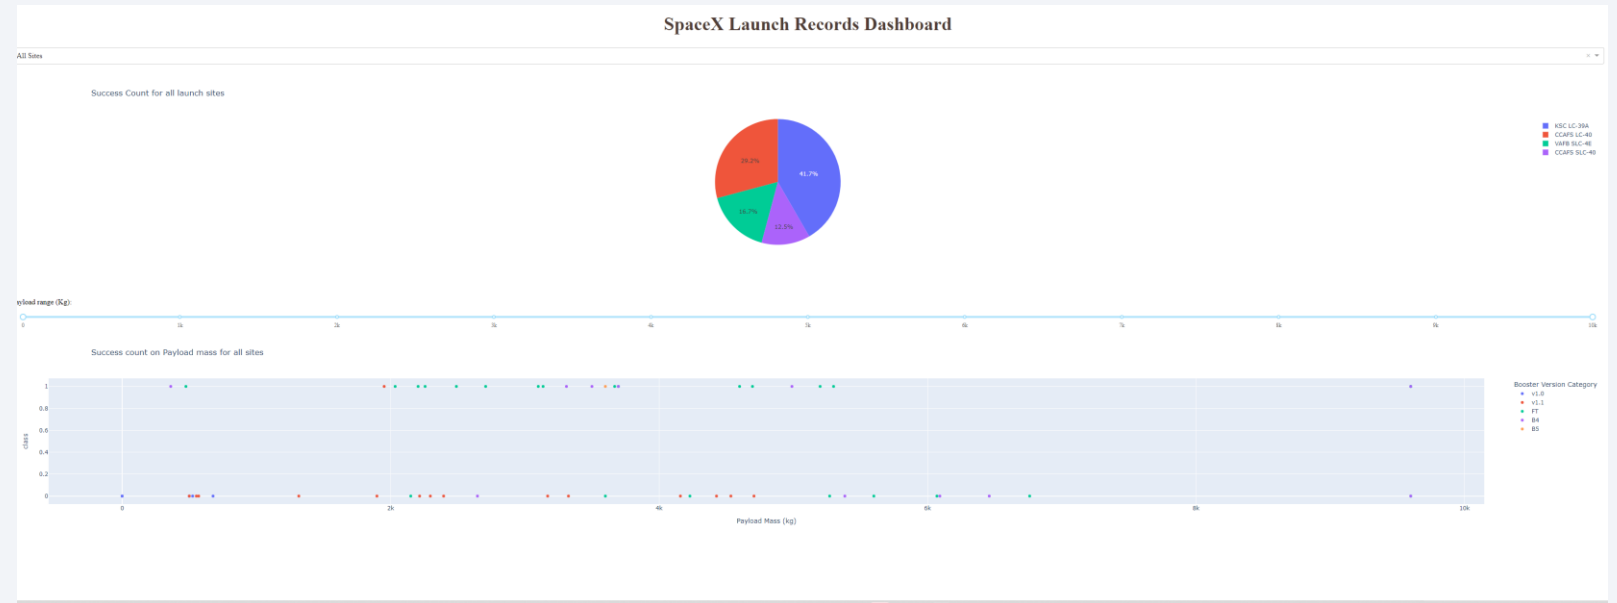
- 1) Launch Sites
- 2) Successful/failed launches
- 3) Proximity of Launch Site

Link to Notebook: [https://jp-tok.dataplatform.cloud.ibm.com/analytics/notebooks/v2/766a1236-b2cc-4fa2-96a4-66bc4edc2317/view?access\\_token=265d45164af6eb93c1733fd68937fb0f009130132302775464ce2c5c4b469c6e](https://jp-tok.dataplatform.cloud.ibm.com/analytics/notebooks/v2/766a1236-b2cc-4fa2-96a4-66bc4edc2317/view?access_token=265d45164af6eb93c1733fd68937fb0f009130132302775464ce2c5c4b469c6e)

# Build a Dashboard with Plotly Dash

Elements Added:

- 1) Dropdown Menu
- 2) Bar Charts
- 3) Scatter Plots



Link to Notebook: <https://github.com/Friedie-FME/Applied-Data-Science-Capstone/blob/main/Dashboard.ipynb>



# Predictive Analysis (Classification)

---

```
In [84]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)
```

```
In [86]: parameters = {'C':[0.01,0.1,1],  
                      'penalty':['l2'],  
                      'solver':['lbfgs']}
```

```
In [87]: parameters = {"C": [0.01, 0.1, 1], 'penalty': ['l2'], 'solver': ['lbfgs']} # l1 Lasso l2 ridge  
lr = LogisticRegression()  
  
logreg_cv = GridSearchCV(lr, parameters, cv=10)  
logreg_cv.fit(X_train, Y_train)
```


Besides **logistics regression**, the following was also run: **Support Vector Machine, Decision Tree, K Nearest Neighbors**

```
In [108... print('Logistics Regression:', logreg_cv.score(X_test, Y_test))  
print('Support Vector Machine:', svm_cv.score(X_test, Y_test))  
print('Decision tree:', tree_cv.score(X_test, Y_test))  
print('K Nearest Neighbors:', knn_cv.score(X_test, Y_test))
```

```
Logistics Regression: 0.8333333333333334  
Support Vector Machine: 0.8333333333333334  
Decision tree: 0.8333333333333334  
K Nearest Neighbors: 0.8333333333333334
```

# Results

- The SVM, KNN and Logistics Regression models have the best prediction accuracy
- KSC LC-39A has the best success rate



```
In [108... print('Logistics Regression:', logreg_cv.score(X_test, Y_test))
print( 'Support Vector Machine:', svm_cv.score(X_test, Y_test))
print('Decision tree:', tree_cv.score(X_test, Y_test))
print('K Nearest Neighbors:', knn_cv.score(X_test, Y_test))

Logistics Regression: 0.8333333333333334
Support Vector Machine: 0.8333333333333334
Decision tree: 0.8333333333333334
K Nearest Neighbors: 0.8333333333333334
```



The background of the slide is an abstract composition. It features a dark blue base color. Overlaid on this are numerous diagonal streaks in shades of red and cyan. A faint, light blue grid pattern is also visible, particularly in the lower-left quadrant. The overall effect is dynamic and technological.

Section 2

# Insights drawn from EDA




# Flight Number vs. Launch Site

## TASK 1: Visualize the relationship between Flight Number and Launch Site

Use the function `catplot` to plot `FlightNumber` vs `LaunchSite`, set the parameter `x` parameter to `FlightNumber`, set the `y` to `Launch Site` and set the parameter `hue` to `'class'`

```
In [4]: # Plot a scatter point chart with x axis to be Flight Number and y axis to be the Launch site, and hue to be the class value
sns.catplot(y="LaunchSite", x="FlightNumber", hue="Class", data=df, aspect = 5)
plt.xlabel("Flight Number",fontsize=20)
plt.ylabel("Launch Site",fontsize=20)
plt.show()
```




# Payload vs. Launch Site

## TASK 2: Visualize the relationship between Payload and Launch Site

We also want to observe if there is any relationship between launch sites and their payload mass.

```
In [5]: # Plot a scatter point chart with x axis to be Pay Load Mass (kg) and y axis to be the Launch site, and hue to be the class value
sns.catplot(y="LaunchSite", x="PayloadMass", hue="Class", data=df, aspect = 5)
plt.xlabel("Launch Site",fontsize=20)
plt.ylabel("Pay load Mass (kg)",fontsize=20)
plt.show()
```



# Success Rate vs. Orbit Type


## TASK 3: Visualize the relationship between success rate of each orbit type

Next, we want to visually check if there are any relationship between success rate and orbit type.

Let's create a `bar chart` for the success rate of each orbit

```
In [6]: # HINT use groupby method on Orbit column and get the mean of Class column
success_rate = df.groupby('Orbit').mean()
success_rate.reset_index(inplace=True)
sns.barplot(x="Orbit", y="Class", data=success_rate, hue='Class')
```

```
Out[6]: <AxesSubplot:xlabel='Orbit', ylabel='Class'>
```



Analyze the plotted bar chart try to find which orbits have high success rate.




# Flight Number vs. Orbit Type

## TASK 4: Visualize the relationship between FlightNumber and Orbit type

For each orbit, we want to see if there is any relationship between FlightNumber and Orbit type.

```
In [7]: # Plot a scatter point chart with x axis to be FlightNumber and y axis to be the Orbit, and hue to be the class value
sns.catplot(y="Orbit", x="FlightNumber", hue="Class", data=df, aspect = 5)
plt.xlabel("Flight Number",fontsize=20)
plt.ylabel("Orbit",fontsize=20)
plt.show()
```




You should see that in the LEO orbit the Success appears related to the number of flights; on the other hand, there seems to be no relationship between flight number when in GTO orbit.

# Payload vs. Orbit Type

## TASK 5: Visualize the relationship between Payload and Orbit type

Similarly, we can plot the Payload vs. Orbit scatter point charts to reveal the relationship between Payload and Orbit type

```
In [10]: # Plot a scatter point chart with x axis to be Payload and y axis to be the Orbit, and hue to be the class value
sns.catplot(y="Orbit", x="PayloadMass", hue="Class", data=df, aspect = 5)
plt.xlabel("Payload",fontsize=20)
plt.ylabel("Orbit",fontsize=20)
plt.show()
```



With heavy payloads the successful landing or positive landing rate are more for Polar, LEO and ISS.

However for GTO we cannot distinguish this well as both positive landing rate and negative landing (unsuccessful mission) are both there here.

# Launch Success Yearly Trend

---

## TASK 6: Visualize the launch success yearly trend

You can plot a line chart with x axis to be `Year` and y axis to be average success rate, to get the average launch success trend.

The function will help you get the year from the date:

```
In [11]: # A function to Extract years from the date
year=[]
def Extract_year(date):
    for i in df["Date"]:
        year.append(i.split("-")[0])
    return year
```

```
In [17]: # Plot a line chart with x axis to be the extracted year and y axis to be the success rate
sns.lineplot(y="successrate", x="year", hue="Class", data=df, aspect = 5)
plt.xlabel("year",fontsize=20)
plt.ylabel("success rate",fontsize=20)
plt.show()
```

# All Launch Site Names

---

## Task 1

Display the names of the unique launch sites in the space mission

In [19]: `%sql SELECT UNIQUE(LAUNCH_SITE) FROM SPACEXDATASET`

\* ibm\_db\_sa://tjg77476:\*\*\*@b0aebb68-94fa-46ec-a1fc-1c999edb6187.c3n41cmd0nqnrk39u98g.databases.appdomain.cloud:31249/bludb  
Done.

Out[19]: **launch\_site**

CCAFS LC-40

CCAFS SLC-40

KSC LC-39A

VAFB SLC-4E

# Launch Site Names Begin with 'CCA'

---

## Task 2

Display 5 records where launch sites begin with the string 'CCA'

```
In [20]: %sql select LAUNCH_SITE from SPACEXDATASET where LAUNCH_SITE LIKE 'CCA%' LIMIT 5
```

```
* ibm_db_sa://tjg77476:***@b0aebb68-94fa-46ec-a1fc-1c999edb6187.c3n41cmd0nqnrk39u98g.databases.appdomain.cloud:31249/bludb  
Done.
```

```
Out[20]: launch_site
```

```
CCAFS LC-40
```

```
CCAFS LC-40
```

```
CCAFS LC-40
```

```
CCAFS LC-40
```

```
CCAFS LC-40
```

# Total Payload Mass

---

## Task 3

Display the total payload mass carried by boosters launched by NASA (CRS)

```
In [21]: %sql SELECT SUM(PAYLOAD_MASS__KG_) from SPACEXDATASET WHERE PAYLOAD LIKE '%CRS%'
```

```
* ibm_db_sa://tjg77476:***@b0aebb68-94fa-46ec-a1fc-1c999edb6187.c3n41cmd0nqnrk39u98g.databases.appdomain.cloud:31249/bludb  
Done.
```

```
Out[21]: 1
```

```
56479
```



# Average Payload Mass by F9 v1.1

---

## Task 4

Display average payload mass carried by booster version F9 v1.1

```
In [22]: %sql SELECT AVG(PAYLOAD_MASS__KG_) from SPACEXDATASET WHERE BOOSTER_VERSION LIKE '%F9 v1.1%'
* ibm_db_sa://tjg77476:***@b0aebb68-94fa-46ec-a1fc-1c999edb6187.c3n41cmd0nqnrk39u98g.databases.appdomain.cloud:31249/bludb
Done.
Out[22]: 1
3226
```

# First Successful Ground Landing Date

---

## Task 5

List the date when the first successful landing outcome in ground pad was achieved.

*Hint: Use min function*

```
In [23]: %sql SELECT MIN(DATE) FROM SPACEXDATASET WHERE MISSION_OUTCOME LIKE 'Success'
```

```
* ibm_db_sa://tjg77476:***@b0aebb68-94fa-46ec-a1fc-1c999edb6187.c3n41cmd0nqnrk39u98g.databases.appdomain.cloud:31249/bludb  
Done.
```

```
Out[23]: 1
```

```
2010-04-06
```

# Successful Drone Ship Landing with Payload between 4000 and 6000

---

## Task 6

List the names of the boosters which have success in drone ship and have payload mass greater than 4000 but less than 6000

```
In [24]: %sql SELECT BOOSTER_VERSION FROM SPACEXDATASET WHERE LANDING__OUTCOME LIKE '%Success (drone ship)%' AND PAYLOAD_MASS__KG_ BETWEEN 4000 AND 6000
```

```
* ibm_db_sa://tjg77476:***@b0aebb68-94fa-46ec-a1fc-1c999edb6187.c3n41cmd0nqnrk39u98g.databases.appdomain.cloud:31249/bludb
Done.
```

```
Out[24]: booster_version
```

```
F9 FT B1022
```

```
F9 FT B1031.2
```

# Total Number of Successful and Failure Mission Outcomes

---

## Task 7

List the total number of successful and failure mission outcomes

```
In [25]: %sql SELECT COUNT(*) FROM SPACEXDATASET GROUP BY MISSION_OUTCOME
```

```
* ibm_db_sa://tjg77476:***@b0aebb68-94fa-46ec-a1fc-1c999edb6187.c3n41cmd0nqnrk39u98g.databases.appdomain.cloud:31249/bludb  
Done.
```

```
Out[25]: 1  
         44  
         1
```

# Boosters Carried Maximum Payload

---

## Task 8

List the names of the booster\_versions which have carried the maximum payload mass. Use a subquery

```
In [26]: %sql SELECT BOOSTER_VERSION FROM SPACEXDATASET WHERE PAYLOAD_MASS__KG_ = (SELECT MAX(PAYLOAD_MASS__KG_) FROM SPACEXDATASET)
* ibm_db_sa://tjg77476:***@b0aebb68-94fa-46ec-a1fc-1c999edb6187.c3n41cmd0nqnrk39u98g.databases.appdomain.cloud:31249/bludb
Done.
```

```
Out[26]: booster_version
```

```
F9 B5 B1048.4
```

```
F9 B5 B1049.4
```

```
F9 B5 B1049.5
```

```
F9 B5 B1060.2
```

```
F9 B5 B1058.3
```

# 2015 Launch Records

---

## Task 9

List the failed landing\_outcomes in drone ship, their booster versions, and launch site names for in year 2015

In [29]: `%sql SELECT LANDING__OUTCOME, BOOSTER_VERSION, LAUNCH_SITE FROM SPACEXDATASET WHERE LANDING__OUTCOME LIKE 'Failure (drone ship)' AND DATE LIKE '%2015%'`

`* ibm_db_sa://tjg77476:***@b0aebb68-94fa-46ec-a1fc-1c999edb6187.c3n41cmd0nqnrk39u98g.databases.appdomain.cloud:31249/bludb`  
Done.

Out[29]: 

landing_outcome	booster_version	launch_site
-----------------	-----------------	-------------

Failure (drone ship)	F9 v1.1 B1012	CCAFS LC-40
----------------------	---------------	-------------



# Rank Landing Outcomes Between 2010-06-04 and 2017-03-20

---

## Task 10

Rank the count of landing outcomes (such as Failure (drone ship) or Success (ground pad)) between the date 2010-06-04 and 2017-03-20, in descending order

```
In [30]: %sql SELECT LANDING__OUTCOME, COUNT(*) AS counts FROM SPACEXDATASET WHERE date BETWEEN '2010-06-04' AND '2017-03-20' GROUP BY LANDING__OUTCOME ORDER BY counts DESC
```

```
* ibm_db_sa://tjg77476:***@b0aebb68-94fa-46ec-a1fc-1c999edb6187.c3n41cmd0nqnrk39u98g.databases.appdomain.cloud:31249/bludb
Done.
```

```
Out[30]:
```

landing_outcome	counts
No attempt	7
Failure (drone ship)	2
Success (drone ship)	2
Success (ground pad)	2
Controlled (ocean)	1
Failure (parachute)	1

A satellite view of Earth from space, showing the curvature of the planet and the glowing city lights of the Eastern United States and parts of Canada at night. The background is a deep blue gradient.

Section 3

# Launch Sites Proximities Analysis

# Location Map of Launch Sites

```
In [8]: # Initial the map
site_map = folium.Map(location=nasa_coordinate, zoom_start=5)

# For each Launch site, add a Circle object based on its coordinate (Lat, Long) values. In addition, add Launch site name as a popup Label
for i, launch_site in launch_sites_df.iterrows():
    coord = [launch_site['Lat'], launch_site['Long']]
    site_name = launch_site['Launch Site']

    circle = folium.Circle(coord, radius=1000, color='#0054d3', fill=True).add_child(folium.Popup(launch_site))
    marker = folium.map.Marker(coord, icon=DivIcon(icon_size=(20,20), icon_anchor=(0,0), html='<div style="font-size: 12; color:#d35400;"><b>%s</b></div>' % launch_site['Launch Site'], ))

    site_map.add_child(circle)
    site_map.add_child(marker)

site_map
```

Out[8]:



# Launch Outcomes on Map

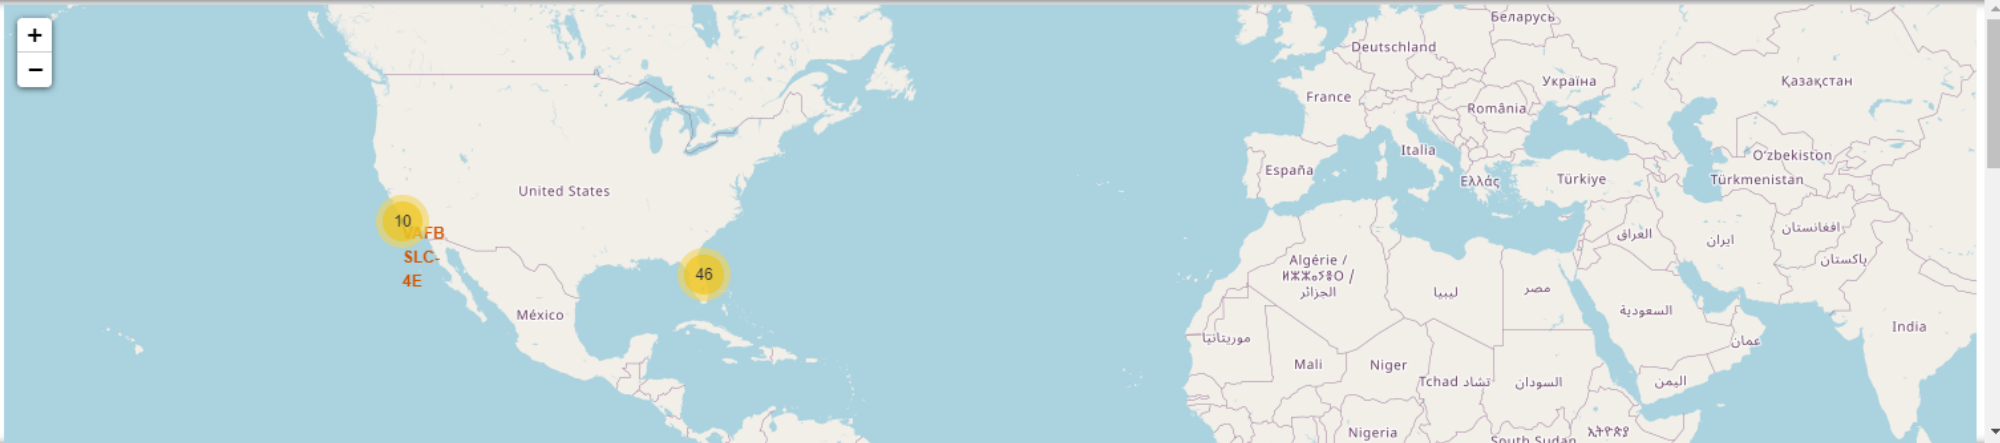
TODO: For each launch result in `spacex_df` data frame, add a `folium.Marker` to `marker_cluster`

```
In [13]: # Add marker_cluster to current site_map
site_map.add_child(marker_cluster)

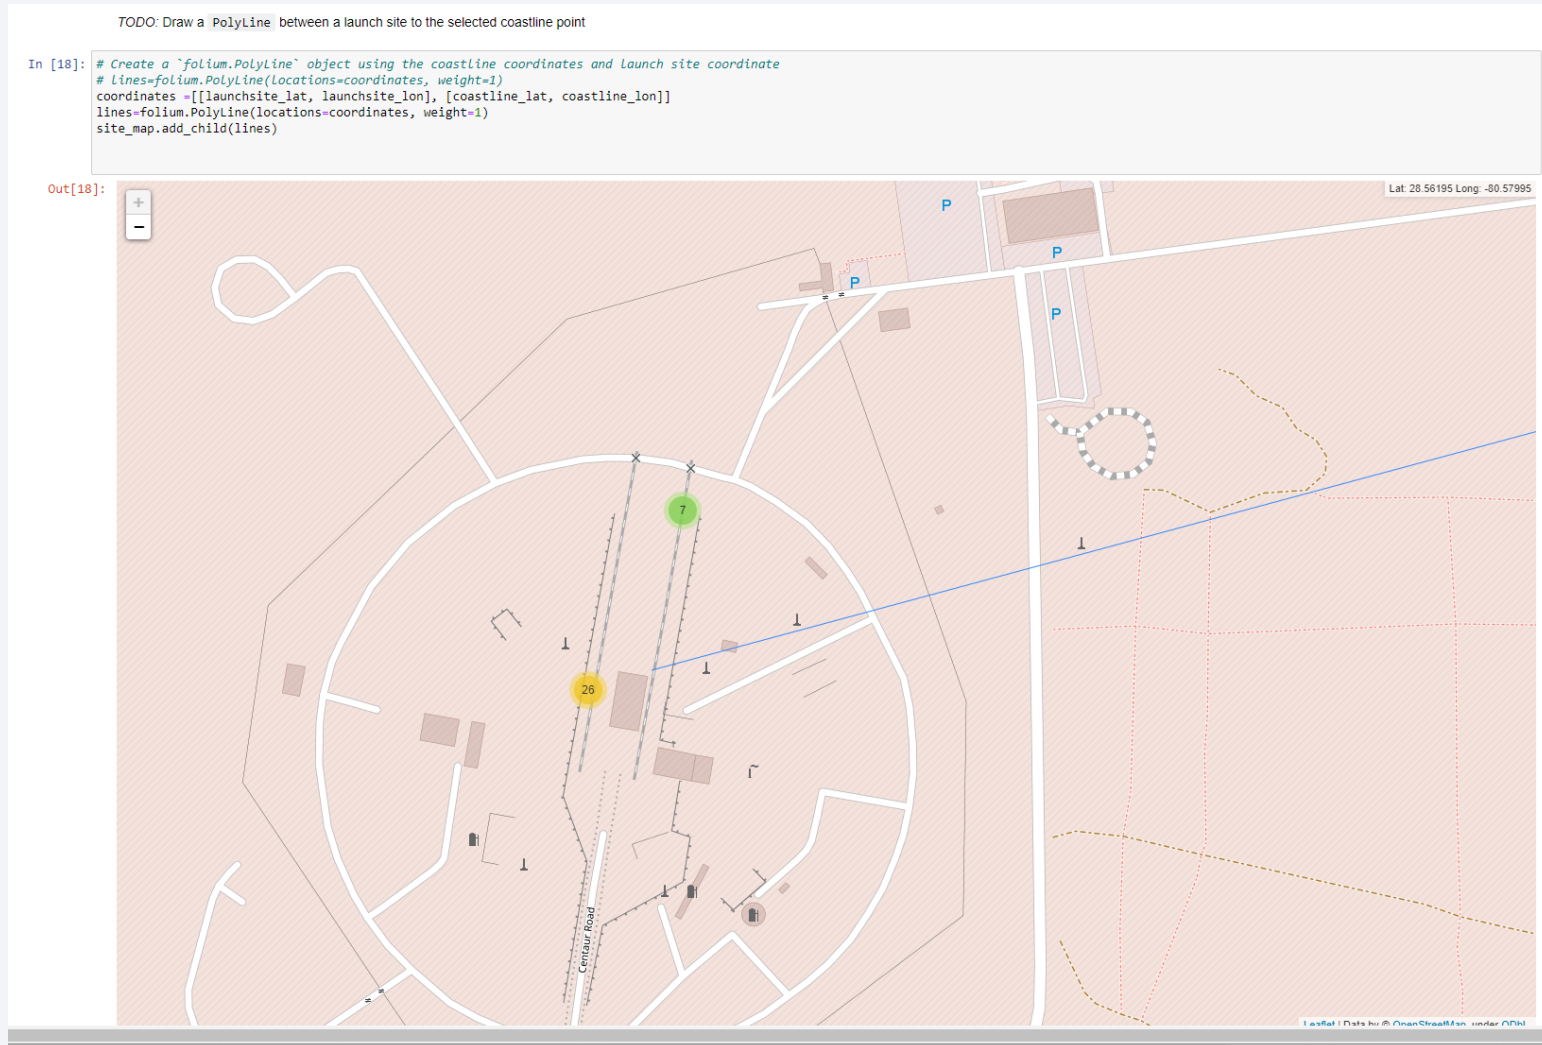
# for each row in spacex_df data frame
# create a Marker object with its coordinate
# and customize the Marker's icon property to indicate if this launch was succeeded or failed,
# e.g., icon=folium.Icon(color='white', icon_color=row['marker_color'])
for index, record in spacex_df.iterrows():
    # TODO: Create and add a Marker cluster to the site map
    marker = folium.Marker([record['Lat'], record['Long']],
                           icon=folium.Icon(color='green', icon_color=record['marker_color']))
    # marker = folium.Marker(...)
    marker_cluster.add_child(marker)

site_map
```

Out[13]:



# Launch Site Proximities on Map







Section 4

# Build a Dashboard with Plotly Dash

# Launch Success for All Sites


---

## SpaceX Launch Records Dashboard

All Sites

×

Success Count for all launch sites



■ KSC LC-39A  
■ CAFS LC-40  
■ VAFB SLC-4E  
■ CAFS SLC-40



# Launch Site with Highest Launch Success


---

## SpaceX Launch Records Dashboard

KSC LC-39A

× ▼


Total Success Launches for site KSC LC-39A



■ 1  
■ 0

Payload range (Kg):

# Payload vs Launch Outcome Scatter Plot



Section 5

# Predictive Analysis (Classification)

# Classification Accuracy


---

```
In [108... print('Logistics Regression:', logreg_cv.score(X_test, Y_test))
print('Support Vector Machine:', svm_cv.score(X_test, Y_test))
print('Decision tree:', tree_cv.score(X_test, Y_test))
print('K Nearest Neighbors:', knn_cv.score(X_test, Y_test))

Logistics Regression: 0.8333333333333334
Support Vector Machine: 0.8333333333333334
Decision tree: 0.8333333333333334
K Nearest Neighbors: 0.8333333333333334
```

# Confusion Matrix

---



# Conclusions

---

Most Interesting Take-Aways:

- The SVM, KNN and Logistics Regression models have the best prediction accuracy
- KSC LC-39A has the best success rate

Thank you!

